2-day Training Course on GPU-enabled Neural Networks
(5 May 2017 and 8 May 2017)

Deep Learning on GPUs is currently boosting AI's massive leap into real-life applications: autopilots, intelligent automated assistants, real-time translation, image recognition, data sequencing and clustering. With the unprecedented computing power of NVIDIA GPUs, many automotive, robotics and big data companies are creating products and services based on a new class of intelligent machines.

This training course is intended for developers willing to rapidly get NVIDIA-based AI technology into new and existing software solutions. In 2 days we will walk from the necessary theory of Recurrent and Convolutional Neural Networks to practical recipes of using Tensorflow and Caffe frameworks.

Hands-ons: All discussed topics will be accompanied by practical sessions, using Caffe and Tensorflow frameworks with NVIDIA backend. Exercises will be conducted on the provided GPU server with NVIDIA Tesla P100 GPUs.

All corresponding presentations and code samples will be available to attendees from the beginning of each training day.

Applied Parallel Computing LLC has been delivering GPU training courses since 2009. Several dozens of courses have been organized all over Europe, both for commercial and academic customers. We work in close partnership with NVIDIA, CUDA Centers of Excellence and Tesla Preferred Partners. In addition to trainings, our company provides GPU porting/optimization services and CUDA certification.

Day 1: Introduction to Machine learning and Neural Networks

Morning (10:00-13:30)

10:00-11:30: lecture

- An overview of GPU performance in various applications
- Machine learning, image recognition, human recognition abilities
- Image filtering
- Image classification patterns, drawbacks
- Bayes classifier (binary classes), classification tasks, examples
- Regression, approximation
- Neural networks, 13th Hilbert problem
- Artificial neuron, weights, activation functions. Perceptron. Layers. Learning approaches. Backpropagation
- Problems of image classification, objects detection, scene understanding explained
- Overview of natural language processing, speech recognition and translation
- Big data, performance
- CPU vs GPU, NVIDIA Kepler, Maxwell and Pascal architecture & performance
- Automotive: Jetson, Parker, Arduino, Raspberry Pi, Robot Operation System, NVIDIA SDK
- Neural networks, supervised and unsupervised learning

- Testing and training sets. Verification, batches
- Self-Organizing Maps

11:30-13:30: hands-on session

- Hands-on, binary classifier, neural network backpropagation, function approximation problem. (CPU)

Afternoon (14:30-18:00)

14:30-16:00: lecture

- Caffe framework. Overview, installation
- Blobs, Layers and Nets. Forward and Backward propagation
- Types of layers: Vision, Loss, Activation /Neuron Layers. Data and Common layers
- Model Format. Prototxt, binaryproto - .caffemodel files
- DIGITS - NVIDIA Deep Learning GPU Training System

16:00-18:00: lecture

- Overview of Google TensorFlow
- Computation graphs. Building graphs, Sessions
- Tensors (ranks, shapes, types), variables
- tf.train.Optimizer – base class for optimizers
- TensorBoard – suite of visualization tools
- Distributed sessions. Creating cluster, distributing training between devices
- Converting data into tensors
- Natural language processing with Recurrent Neural Networks

Day 2: Convolutional Neural Networks and Caffe

Morning (10:00-13:30)

10:00-11:30: lecture

- Solvers and methods for neural network training
- Convolutional networks, structures, parameters
- LeNet, AlexNet, GoogleNet, VGG
- Guidelines on choosing neural network structure
- Open-source datasets
- Detectnet – Detecting objects
- NVIDIA TensorRT– optimizing and deploying trained network
- Combining CAFFE, Tensorflow and custom C/C++ code

11:30-13:30: hands-on session

- DC motor transient model approximation
- Equalizing and synchronizing DC motors

Afternoon (14:30-18:00)

14:30-16:00: hands-on session

- Recognizing robot and calculating it's coordinates on image. Detectnet
- Recognizing other objects on scene. Detectnet

16:00-18:00: hands-on session

- Understanding commands using RNN (Tensorflow)
- Navigating robot, using neural networks for DC motors controlling,
- Avoiding obstacles, understanding and interpreting commands, recognizing objects
- Live demo