2-day OpenACC training course at Universität Jena

(4-5 December 2017)

CUDA and OpenCL are two programming languages that allow to develop GPU applications with great level of detail, but require advanced knowledge of GPU architecture and extensive code re-writing. Many cross-platform or prototype applications do not actually require deep GPU porting, if they perform data processing using standard algorithms, e.g. reduction. In such cases, fair performance numbers could be achieved with simplified *directive-based* language extensions. OpenACC is de facto an industry standard for C/Fortran directive extensions for porting code to GPUs. This training course walks through OpenACC programming technology, from basics to advanced practices used in real applications. As a case study, the course presents a practical session on implementing OpenACC spectral space solver for a Poisson problem.

Hands-ons: All discussed topics will be accompanied with practical sessions, using PGI OpenACC compiler with NVIDIA backend. Exercises will be conducted either on the provided remote GPU server or on the customer's local system.

All corresponding presentations and code samples will be available to attendees from the beginning of each training day.

Applied Parallel Computing LLC is delivering GPU training courses since 2009. Several dozens of courses have been organized all over Europe, both for commercial and academic customers. We work in close partnership with NVIDIA, CUDA Centers of Excellence and Tesla Preferred Partners. In addition to trainings, our company provides GPU porting/optimization services and CUDA certification.

Day 1: Introduction to OpenACC

Morning (10:00-13:30)

10:00-11:30: lecture

- Advantages of OpenACC.
- Execution model: gangs, workers and vectors – three levels of coarse-grain and fine-grain parallelism. SIMD instructions.
- OpenACC memory model: host and accelerator address spaces.
- OpenACC directive syntax in C and Fortran. Main directives: *parallel* and *kernels* – offloading code regions to accelerator, *loop* – detailed parallelization parameters for each loop.
- Examples of *vector addition* and *reduction* explained.

12:00-13:30: hands-on session

- PGI compiler for NVIDIA GPUs.
- Compilation and deployment of OpenACC code examples.

Afternoon (14:30-18:00)

14:30-16:00: lecture

- Understanding PGI OpenACC compiler output. Compiler flags and environment variables for detailed analysis and performance reports. Compiler limitations in dependencies tracking, enforcing data independence.
- Data clauses: *deviceptr*, *copy*, *copyin*, *copyout*, *create*, *delete*, *present*, *present_or_(-copy, -copyin, -copyout, -create)*.
- Organizing data persistence regions using *data* directives.
- Non-structured data lifetime with *enter data* and *exit data* directives.
- Accessing global variables.
- Additional data management directives: *cache*, *update*, *declare*.
- External dependencies in OpenACC kernels, functions inlining. OpenACC *routine* directive, separate compilation and procedure calls.

16:30-18:00: hands-on session

- "Fill-in" exercise on reimplementing CUFFT + CPU version of 2D Poisson solver into efficient CUFFT + OpenACC version.

Day 2: Advanced OpenACC programming

Morning (08:30-12:00)

08:30-10:00: lecture

- Profiling GPU kernels in OpenACC application. *Time* option, *PGI_ACC_TIME* environment variable. Profiling with *pgprof*.
- Expressing locality in loop nest with *tile* directive.

- Organizing asynchronous execution using *async* clause and *wait* directive.
- *Atomic* directive. Allowed operations. Restrictions.
- Advanced code optimization practices, by examples. Restrictions, common mistakes, workarounds.

10:00-10:20: Coffee break
10:20-11:45: lecture

- GPU-enabled scientific libraries
- OpenACC interoperation with CUDA C, CUDA Fortran and GPU-enabled libraries.

11:45-12:00: Q & A

Prerequisites

- Beamer
- Chrome or Firefox web browser
- (optional) Open (unfirewalled) connection to farm.parallel-computing.pro (188.44.42.33) over port 22
- (optional) PuTTY SSH client