



2-day OpenACC training course at Universität Leipzig

(6-7 December 2017)

CUDA and OpenCL are two programming languages that allow to develop GPU applications with great level of detail, but require advanced knowledge of GPU architecture and extensive code re-writing. Many cross-platform or prototype applications do not actually require deep GPU porting, if they perform data processing using standard algorithms, e.g. reduction. In such cases, fair performance numbers could be achieved with simplified *directive-based* language extensions. OpenACC is de facto an industry standard for C/Fortran directive extensions for porting code to GPUs. This training course walks through OpenACC programming technology, from basics to advanced practices used in real applications. As a case study, the course presents a practical session on using OpenACC for Deep Learning convolutional neural networks.

Hands-ons: All discussed topics will be accompanied with practical sessions, using PGI OpenACC compiler with NVIDIA backend. Exercises will be conducted either on the provided remote GPU server or on the customer's local system.

All corresponding presentations will be available to attendees in printed handouts.

Applied Parallel Computing LLC is delivering GPU training courses since 2009. Several dozens of courses have been organized all over Europe, both for commercial and academic customers. We work in close partnership with NVIDIA, CUDA Centers of Excellence and Tesla Preferred Partners. In addition to trainings, our company provides GPU porting/optimization services and [CUDA certification](#).

Day 1: Introduction to OpenACC

Morning (10:00-13:30)

10:00-11:15: lecture

- An overview of GPU performance in various applications
- Brief intercomparison of different types of accelerators
- Key programming principles to achieve high GPU performance

11:15-11:30: coffee break

11:30-12:30: lecture

- Advantages of OpenACC in comparison to CUDA.
- Execution model: gangs, workers and vectors – three levels of coarse-grain and fine-grain parallelism. SIMD instructions.
- OpenACC memory model: host and accelerator address spaces.
- OpenACC directive syntax in C and Fortran. Main directives: *parallel* and *kernels* – offloading code regions to

accelerator, *loop* – detailed parallelization parameters for each loop.

- Examples of *vector addition* and *reduction* explained, in comparison to CUDA.

12:30-13:30: [hands-on session](#)

- PGI compiler for NVIDIA GPUs. GCC open-source OpenACC compiler.
- Compilation and deployment of OpenACC code examples.

Afternoon (14:30-18:00)

14:30-16:00: lecture

- Understanding PGI OpenACC compiler output. Compiler flags and environment variables for detailed analysis and performance reports. Compiler limitations in dependencies tracking, enforcing data independence.
- Data clauses: *deviceptr*, *copy*, *copyin*, *copyout*, *create*, *delete*, *present*, *present_or_(-copy, -copyin, -copyout, -create)*.

- Organizing data persistence regions using *data* directives.

16:00-16:20: coffee break

16:20-18:00: [hands-on session](#)

- Profiling GPU kernels in OpenACC application. *Time* option, *PGI_ACC_TIME* environment variable. Profiling with *pgprof*.
- “Fill-in” exercise on implementing wave propagation stencil in OpenACC (*wave13pt*). Adding OpenACC directives step by step.

Day 2: Advanced OpenACC programming

Morning (09:00-12:30)

09:00-10:30: lecture

- Non-structured data lifetime with *enter data* and *exit data* directives.
- Additional data management directives: *cache*, *update*, *declare*.
- Expressing locality in loop nest with *tile* directive.
- Organizing asynchronous execution using *async* clause and *wait* directive.
- *Atomic* directive. Allowed operations. Restrictions.

10:30-11:00: coffee break

11:00-12:30: [hands-on session](#)

- Using non-structured data lifetime directives
- Example of OpenACC loop tiling, effect on performance
- Advanced code optimization practices, by examples. Restrictions, common mistakes, workarounds.
- External dependencies in OpenACC kernels, functions inlining. OpenACC *routine* directive, separate compilation and procedure calls.
- Accessing global variables.

Afternoon (13:30-17:00)

13:30-15:00: lecture

- GPU-enabled scientific libraries
- OpenACC interoperation with CUDA and GPU-enabled libraries.
- Introduction into neural networks and deep learning

15:00-15:20: coffee break

15:20-16:30: [hands-on session](#)

- Porting Batched Gradient Descent method for Convolutional Neural Network Deep Learning onto GPU using OpenACC.

16:30-17:00: [Q & A](#)

Prerequisites

- Beamer
- Open (unfirewalled) connection to farm.parallel-computing.pro (188.44.42.33) over port 22
- PuTTY SSH client